

# Exploring Big Data

---

## Project Description

This task is based on a synthesised transaction dataset containing 3 months' worth of transactions for 100 hypothetical customers. It contains purchases, recurring transactions, and salary transactions.

This task will assess some of the core skills expected in a Big Data engineer at ANZ, particularly your familiarity with Apache Spark.

Using each API, perform the following transformation steps using the synthetic transaction file as input referenced as an input argument to your program. Output the results to a local file.

- Project only the records where status=authorized AND card\_present\_flag=0
- Split the long\_lat and merchant\_long\_lat fields into long, lat and merch\_long, merch\_lat fields
- Output the data as a CSV file

## Import Libraries

```
In [1]: import numpy as np
from numpy import count_nonzero, median, mean
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import random

# Plotly
# import plotly.express as px
# import plotly.offline as py
# import plotly.graph_objs as go

import sweetviz

import statsmodels.api as sm
import statsmodels.formula.api as smf
from statsmodels.formula.api import ols
# import researchpy as rp

import datetime
from datetime import datetime, timedelta

# import eli5
# from IPython.display import display

#import os
#import zipfile
import scipy.stats
from collections import Counter

import sklearn
# from sklearn.preprocessing import StandardScaler, MinMaxScaler, LabelEncoder, OneHotEn
# from sklearn.linear_model import LinearRegression, LogisticRegression, ElasticNet, Las
# from sklearn.model_selection import cross_val_score, train_test_split
```

```

# from sklearn.metrics import accuracy_score, auc, classification_report, confusion_matrix
# from sklearn.metrics import plot_confusion_matrix, plot_roc_curve

# from sklearn.linear_model import ElasticNet, Lasso, LinearRegression, LogisticRegression
# from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, ExtraTreeClassifier
# from sklearn.svm import SVC, SVR, LinearSVC, LinearSVR
# from sklearn.naive_bayes import GaussianNB, MultinomialNB

%matplotlib inline
#sets the default autosave frequency in seconds
%autosave 60
sns.set_style('dark')
sns.set(font_scale=1.2)

plt.rc('axes', titlesize=9)
plt.rc('axes', labelszize=14)
plt.rc('xtick', labelszize=12)
plt.rc('ytick', labelszize=12)

import warnings
warnings.filterwarnings('ignore')

# Use Feature-Engine library
#import feature_engine
#from feature_engine import imputation as mdi
#from feature_engine.outlier_removers import Winsorizer
#from feature_engine import categorical_encoders as ce
#from feature_engine.discretisation import EqualWidthDiscretiser, EqualFrequencyDiscretiser
#from feature_engine.discretisation import ArbitraryDiscretiser, DecisionTreeDiscretiser
#from feature_engine.encoding import OrdinalEncoder

pd.set_option('display.max_columns',None)
#pd.set_option('display.max_rows',None)
pd.set_option('display.width', 1000)
pd.set_option('display.float_format', '{:.2f}'.format)

random.seed(0)
np.random.seed(0)
np.set_printoptions(suppress=True)

```

Autosaving every 60 seconds

## Exploratory Data Analysis

In [2]: `df = pd.read_csv("BIGDATAContentTask.csv")`

In [3]: `df.head()`

Out[3]:

	status	card_present_flag	bpay_biller_code	account	currency	long_lat	txn_description	merchant_id
0	authorized	1.00	NaN	ACC-1598451071	AUD	153.41 -27.95	POS	81c48296-73be-44a7-befa-d053f48ce7cd
1	authorized	0.00	NaN	ACC-1598451071	AUD	153.41 -27.95	SALES-POS	830a451c-316e-4a6a-bf25-e37caedca49e
2	authorized	1.00	NaN	ACC-1222300524	AUD	151.23 -33.94	POS	835c231d-8cdf-4e96-859d-e9d571760cf0

3	authorized	1.00	NaN	ACC-1037050564	AUD	153.10 -27.66	SALES-POS	48514682- c78a-4a88- b0da- 2d6302e64673
4	authorized	1.00	NaN	ACC-1598451071	AUD	153.41 -27.95	SALES-POS	b4e02c10- 0852-4273- b8fd- 7b3395e32eb0

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12043 entries, 0 to 12042
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   status                12043 non-null  object
1   card_present_flag     7717 non-null   float64
2   bpay_biller_code      885 non-null    object
3   account               12043 non-null  object
4   currency              12043 non-null  object
5   long_lat              12043 non-null  object
6   txn_description       12043 non-null  object
7   merchant_id           7717 non-null   object
8   merchant_code         883 non-null    float64
9   first_name            12043 non-null  object
10  balance               12043 non-null  float64
11  date                  12043 non-null  object
12  gender                12043 non-null  object
13  age                   12043 non-null  int64
14  merchant_suburb       7717 non-null   object
15  merchant_state        7717 non-null   object
16  extraction            12043 non-null  object
17  amount                12043 non-null  float64
18  transaction_id        12043 non-null  object
19  country               12043 non-null  object
20  customer_id           12043 non-null  object
21  merchant_long_lat     7717 non-null   object
22  movement              12043 non-null  object
dtypes: float64(4), int64(1), object(18)
memory usage: 2.1+ MB
```

```
In [5]: df.describe()
```

Out[5]:

	card_present_flag	merchant_code	balance	age	amount
count	7717.00	883.00	12043.00	12043.00	12043.00
mean	0.80	0.00	14704.20	30.58	187.93
std	0.40	0.00	31503.72	10.05	592.60
min	0.00	0.00	0.24	18.00	0.10
25%	1.00	0.00	3158.59	22.00	16.00
50%	1.00	0.00	6432.01	28.00	29.00
75%	1.00	0.00	12465.94	38.00	53.66
max	1.00	0.00	267128.52	78.00	8835.98

```
In [6]: df.columns
```

Out[6]: Index(['status', 'card\_present\_flag', 'bpay\_biller\_code', 'account', 'currency', 'long\_lat', 'txn\_description', 'merchant\_id', 'merchant\_code', 'first\_name', 'balance', 'date', 'gender', 'age', 'merchant\_suburb', 'merchant\_state', 'extraction', 'amount', 'transaction\_id', 'country', 'customer\_id', 'merchant\_long\_lat', 'movement'], dtype='object')

In [7]: df.status.value\_counts()

Out[7]: authorized 7717  
posted 4326  
Name: status, dtype: int64

In [8]: df2 = df[df.status == "authorized"]  
df2

	status	card_present_flag	bpay_biller_code	account	currency	long_lat	txn_description	merchant_id
0	authorized	1.00	NaN	ACC-1598451071	AUD	153.41-27.95	POS	81c48273be-44b1d053f48ce
1	authorized	0.00	NaN	ACC-1598451071	AUD	153.41-27.95	SALES-POS	830a45316e-4ab1e37caedca
2	authorized	1.00	NaN	ACC-1222300524	AUD	151.23-33.94	POS	835c238cdf-4e85e9d571760
3	authorized	1.00	NaN	ACC-1037050564	AUD	153.10-27.66	SALES-POS	485146c78a-4ab1c2d6302e64
4	authorized	1.00	NaN	ACC-1598451071	AUD	153.41-27.95	SALES-POS	b4e02c0852-42b17b3395e32
...	...	...	...	...	...	...	...	...
12038	authorized	0.00	NaN	ACC-3021093232	AUD	149.83-29.47	POS	32aa73b7c2-41b16271b96ce
12039	authorized	1.00	NaN	ACC-1608363396	AUD	151.22-33.87	SALES-POS	296a058552-4f8ac37065b5
12040	authorized	1.00	NaN	ACC-3827517394	AUD	151.12-33.89	POS	e5975ae08f7-47a324cc0e35e
12041	authorized	1.00	NaN	ACC-2920611728	AUD	144.96-37.76	SALES-POS	af4905591d-4fbc27730b70e
12042	authorized	1.00	NaN	ACC-1443681913	AUD	150.92-33.77	SALES-POS	f31f4b2040-40a1b141bb274

7717 rows × 23 columns

```
In [9]: df3 = df2[df2.card_present_flag == 0.00]  
df3
```

```
Out[9]:
```

	status	card_present_flag	bpay_biller_code	account	currency	long_lat	txn_description	merchant
1	authorized	0.00	NaN	ACC- 1598451071	AUD	153.41 -27.95	SALES-POS	830a45 316e-4a bf e37caedca
21	authorized	0.00	NaN	ACC- 2890243754	AUD	153.32 -27.93	POS	7e8bf6 e724-43 a4 3538a3e27
23	authorized	0.00	NaN	ACC- 2615038700	AUD	145.35 -38.03	POS	354f40 55bc-4a a0 c7faede29
29	authorized	0.00	NaN	ACC- 1710017148	AUD	150.82 -34.01	SALES-POS	4af250 a1a4-46 90 240d790e5
31	authorized	0.00	NaN	ACC- 3485804958	AUD	138.52 -35.01	POS	a08935 99a8-4f b7 f8de51ca0
...	...	...	...	...	...	...	...	...
12012	authorized	0.00	NaN	ACC- 3954677887	AUD	115.72 -32.28	SALES-POS	4995bf 13d8-4f a9 d5331da83
12015	authorized	0.00	NaN	ACC- 1598451071	AUD	153.41 -27.95	POS	e4758c e8d8-49 99 a823a86dc
12016	authorized	0.00	NaN	ACC- 2249586092	AUD	115.98 -32.07	SALES-POS	23eccb 684e-43 b9 dc307517c
12031	authorized	0.00	NaN	ACC- 1443681913	AUD	150.92 -33.77	SALES-POS	6fcdc9 3548-40 a2 9dbc6cb64
12038	authorized	0.00	NaN	ACC- 3021093232	AUD	149.83 -29.47	POS	32aa73 b7c2-41 b1 6271b96ce

1523 rows × 23 columns

```
In [10]: df3.reset_index(inplace=True, drop=True)
```

```
In [11]: df3
```

Out[11]:

	status	card_present_flag	bpay_billir_code	account	currency	long_lat	txn_description	merchant_
0	authorized	0.00	NaN	ACC-1598451071	AUD	153.41 -27.95	SALES-POS	830a451 316e-4a6 bf2 e37caedca4
1	authorized	0.00	NaN	ACC-2890243754	AUD	153.32 -27.93	POS	7e8bf66 e724-435 a40 3538a3e27b
2	authorized	0.00	NaN	ACC-2615038700	AUD	145.35 -38.03	POS	354f40c 55bc-4a8 a00 c7faede29fi
3	authorized	0.00	NaN	ACC-1710017148	AUD	150.82 -34.01	SALES-POS	4af2504 a1a4-468 90b 240d790e53
4	authorized	0.00	NaN	ACC-3485804958	AUD	138.52 -35.01	POS	a08935a 99a8-49f b73 f8de51ca0al
...	...	...	...	...	...	...	...	...
1518	authorized	0.00	NaN	ACC-3954677887	AUD	115.72 -32.28	SALES-POS	4995bfd 13d8-4c a9f d5331da83fi
1519	authorized	0.00	NaN	ACC-1598451071	AUD	153.41 -27.95	POS	e4758c3 e8d8-49b 990 a823a86dca
1520	authorized	0.00	NaN	ACC-2249586092	AUD	115.98 -32.07	SALES-POS	23eccb6 684e-432 b95 dc307517c8
1521	authorized	0.00	NaN	ACC-1443681913	AUD	150.92 -33.77	SALES-POS	6fcdc95 3548-40b a2c 9dbc6cb64el
1522	authorized	0.00	NaN	ACC-3021093232	AUD	149.83 -29.47	POS	32aa73d b7c2-416 b14 6271b96ce7

1523 rows × 23 columns


# Data Visualization

## Univariate Data Exploration

```
In [12]: df3.hist(bins=50, figsize=(20,10))
plt.suptitle('Histogram Feature Distribution', x=0.5, y=1.02, ha='center', fontsize=20)
```

```
plt.tight_layout()
plt.show()
```


Histogram Feature Distribution



```
In [13]: df3.boxplot(figsize=(20,10))
plt.suptitle('BoxPlots Feature Distribution', x=0.5, y=1.02, ha='center', fontsize=20)

plt.tight_layout()
plt.show()
```

BoxPlots Feature Distribution



```
In [14]: df3.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1523 entries, 0 to 1522
Data columns (total 23 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   status                1523 non-null   object  
 1   card_present_flag     1523 non-null   float64 
 2   bpay_biller_code      0 non-null      object
```

```
3   account      1523 non-null object
4   currency     1523 non-null object
5   long_lat     1523 non-null object
6   txn_description 1523 non-null object
7   merchant_id  1523 non-null object
8   merchant_code 0 non-null float64
9   first_name   1523 non-null object
10  balance      1523 non-null float64
11  date         1523 non-null object
12  gender       1523 non-null object
13  age          1523 non-null int64
14  merchant_suburb 1523 non-null object
15  merchant_state 1523 non-null object
16  extraction   1523 non-null object
17  amount       1523 non-null float64
18  transaction_id 1523 non-null object
19  country      1523 non-null object
20  customer_id  1523 non-null object
21  merchant_long_lat 1523 non-null object
22  movement     1523 non-null object
dtypes: float64(4), int64(1), object(18)
memory usage: 273.8+ KB
```

## Data Preprocessing

### Feature Engineering

```
In [15]: df3.columns
```

```
Out[15]: Index(['status', 'card_present_flag', 'bpay_biller_code', 'account', 'currency', 'long_lat', 'txn_description', 'merchant_id', 'merchant_code', 'first_name', 'balance', 'date', 'gender', 'age', 'merchant_suburb', 'merchant_state', 'extraction', 'amount', 'transaction_id', 'country', 'customer_id', 'merchant_long_lat', 'movement'], dtype='object')
```

```
In [16]: df3.drop(['status', 'card_present_flag', 'bpay_biller_code', 'account', 'currency', 'long_lat', 'txn_description', 'merchant_id', 'merchant_code', 'first_name'], inplace=True)
```

```
In [17]: df3.head()
```

```
Out[17]:
```

	balance	date	gender	age	merchant_suburb	merchant_state	extraction	amount	
0	21.20	1/8/2018	F	26	Sydney	NSW	2018-08-01T01:13:45.000+0000	14.19	13270a2a90
1	275.93	1/8/2018	M	37	Lismore	NSW	2018-08-01T08:19:14.000+0000	24.77	1f12467d33
2	30583.15	1/8/2018	F	43	Mordialloc	VIC	2018-08-01T08:47:48.000+0000	12.08	49417bad3f
3	1625.34	1/8/2018	F	19	Alexandria	NSW	2018-08-01T09:11:00.000+0000	11.57	82acf0379
4	12529.59	1/8/2018	F	34	Findon	SA	2018-08-01T09:19:06.000+0000	33.89	89050ee5c5

```
In [18]: df3.drop(['date', 'merchant_suburb', 'extraction', 'transaction_id', 'country', 'customer_id', 'merchant_long_lat', 'movement'], inplace=True)
```

```
In [19]: df3.head()
```

```
Out[19]:
```

	balance	gender	age	merchant_state	amount	movement
--	---------	--------	-----	----------------	--------	----------



<b>0</b>	21.20	F	26	NSW	14.19	debit
<b>1</b>	275.93	M	37	NSW	24.77	debit
<b>2</b>	30583.15	F	43	VIC	12.08	debit
<b>3</b>	1625.34	F	19	NSW	11.57	debit
<b>4</b>	12529.59	F	34	SA	33.89	debit

In [ ]:

## Save to CSV

In [20]: `df3.to_csv("final.csv", index=False)`

## Regression Analysis

### Logistic Regression (StatsModel)

In [21]: `df3.columns`

Out[21]: `Index(['balance', 'gender', 'age', 'merchant_state', 'amount', 'movement'], dtype='object')`

In [22]: `df3.movement.value_counts()`

Out[22]: `debit 1523  
Name: movement, dtype: int64`

In [23]: `df4 = pd.get_dummies(data=df3)`

In [24]: `df4`

Out[24]:

	balance	age	amount	gender_F	gender_M	merchant_state_ACT	merchant_state_NSW	merchant_state_NT
<b>0</b>	21.20	26	14.19	1	0	0	1	0
<b>1</b>	275.93	37	24.77	0	1	0	1	0
<b>2</b>	30583.15	43	12.08	1	0	0	0	0
<b>3</b>	1625.34	19	11.57	1	0	0	1	0
<b>4</b>	12529.59	34	33.89	1	0	0	0	0
...	...	...	...	...	...	...	...	...
<b>1518</b>	9901.03	47	15.91	1	0	0	0	0
<b>1519</b>	2194.26	26	25.88	1	0	0	0	0
<b>1520</b>	12963.75	19	9.90	0	1	0	0	0
<b>1521</b>	5540.27	31	70.51	0	1	0	1	0
<b>1522</b>	14054.14	30	9.79	1	0	0	0	0

1523 rows × 14 columns

In [25]: `df4.columns`

```
Out[25]: Index(['balance', 'age', 'amount', 'gender_F', 'gender_M', 'merchant_state_ACT', 'merchant_state_NSW', 'merchant_state_NT', 'merchant_state QLD', 'merchant_state_SA', 'merchant_state_TAS', 'merchant_state_VIC', 'merchant_state_WA', 'movement_debit'], dtype='object')
```

```
In [26]: y = df4['movement_debit']  
X = df4[['balance', 'age', 'amount', 'gender_F', 'gender_M', 'merchant_state_ACT', 'merc
```

```
In [27]: X = sm.add_constant(X)
```

```
In [28]: model = sm.Logit(y, X).fit()
```

-----  
**PerfectSeparationError**

Traceback (most recent call last)

Input In [28], in <cell line: 1>()

----> 1 model = sm.Logit(y, X).fit()

File C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\discrete\discrete\_model.py:1983, in Logit.fit(self, start\_params, method, maxiter, full\_output, disp, callback, \*\*kwargs)

```
1980 @Appender(DiscreteModel.fit.__doc__)  
1981 def fit(self, start_params=None, method='newton', maxiter=35,  
1982         full_output=1, disp=1, callback=None, **kwargs):  
-> 1983     bnryfit = super().fit(start_params=start_params,  
1984                           method=method,  
1985                           maxiter=maxiter,  
1986                           full_output=full_output,  
1987                           disp=disp,  
1988                           callback=callback,  
1989                           **kwargs)  
1991     discretefit = LogitResults(self, bnryfit)  
1992     return BinaryResultsWrapper(discretefit)
```

File C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\discrete\discrete\_model.py:230, in DiscreteModel.fit(self, start\_params, method, maxiter, full\_output, disp, callback, \*\*kwargs)

```
227 else:  
228     pass # TODO: make a function factory to have multiple call-backs  
-> 230 mlefit = super().fit(start_params=start_params,  
231                       method=method,  
232                       maxiter=maxiter,  
233                       full_output=full_output,  
234                       disp=disp,  
235                       callback=callback,  
236                       **kwargs)  
238 return mlefit
```

File C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\base\model.py:563, in LikelihoodModel.fit(self, start\_params, method, maxiter, full\_output, disp, fargs, callback, retall, skip\_hessian, \*\*kwargs)

```
560     del kwargs["use_t"]  
562     optimizer = Optimizer()  
-> 563     xopt, retvals, optim_settings = optimizer.fit(f, score, start_params,  
564                                                  fargs, kwargs,  
565                                                  hessian=hess,  
566                                                  method=method,  
567                                                  disp=disp,  
568                                                  maxiter=maxiter,  
569                                                  callback=callback,  
570                                                  retall=retall,  
571                                                  full_output=full_output)  
572 # Restore cov_type, cov_kwds and use_t  
573 optim_settings.update(kwds)
```

File C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\base\optimizer.py:241, in Op

```

optimizer._fit(self, objective, gradient, start_params, fargs, kwargs, hessian, method, ma
xiter, full_output, disp, callback, retall)
238     fit_funcs.update(extra_fit_funcs)
240     func = fit_funcs[method]
--> 241     xopt, retvals = func(objective, gradient, start_params, fargs, kwargs,
242                             disp=disp, maxiter=maxiter, callback=callback,
243                             retall=retall, full_output=full_output,
244                             hess=hessian)
246     optim_settings = {'optimizer': method, 'start_params': start_params,
247                       'maxiter': maxiter, 'full_output': full_output,
248                       'disp': disp, 'fargs': fargs, 'callback': callback,
249                       'retall': retall, "extra_fit_funcs": extra_fit_funcs}
250     optim_settings.update(kwargs)

File C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\base\optimizer.py:443, in _f
it_newton(f, score, start_params, fargs, kwargs, disp, maxiter, callback, retall, full_o
utput, hess, ridge_factor)
441         history.append(newparams)
442         if callback is not None:
--> 443             callback(newparams)
444         iterations += 1
445     fval = f(newparams, *fargs) # this is the negative likelihood

File C:\ProgramData\Anaconda3\lib\site-packages\statsmodels\discrete\discrete_model.py:2
14, in DiscreteModel._check_perfect_pred(self, params, *args)
211     if (self.raise_on_perfect_prediction and
212         np.allclose(fittedvalues - endog, 0)):
213         msg = "Perfect separation detected, results not available"
--> 214         raise PerfectSeparationError(msg)

PerfectSeparationError: Perfect separation detected, results not available

```

```
In [ ]: model.summary()
```

```
In [ ]: logitfit = smf.logit(formula = 'DF ~ Debt_Service_Coverage + cash_security_to_curLiab +
```

```
In [ ]: logitfit = smf.logit(formula = 'DF ~ TNW + C(seg2)', data = hgcdev).fit()
```

```
In [ ]:
```